

# FNLP Lecture 9: Algorithms for HMMs

Henry S. Thompson

Based on slides by Sharon Goldwater

12 February 2019

# Recap: HMM

- Elements of HMM:
  - Set of states (tags)
  - Output alphabet (word types)
  - Start state (beginning of sentence)
  - State transition probabilities
  - Output probabilities from each state

# More general notation

- Previous lecture:
  - Sequence of tags  $T = t_1 \dots t_n$
  - Sequence of words  $S = w_1 \dots w_n$
- This lecture:
  - Sequence of states  $Q = q_1 \dots q_T$
  - Sequence of outputs  $O = o_1 \dots o_T$
  - So  $t$  is now a time step, not a tag! And  $T$  is the sequence length.

# Recap: HMM

- Given a sentence  $O = o_1 \dots o_T$  with tags  $Q = q_1 \dots q_T$ , compute  $P(O, Q)$  as:

$$P(O, Q) = \prod_{t=1}^T P(o_t | q_t) P(q_t | q_{t-1})$$

- But we want to find  $\operatorname{argmax}_Q P(Q|O)$  without enumerating all possible  $Q$ 
  - Use Viterbi algorithm to store partial computations.

# Today's lecture

- What algorithms can we use to
  - Efficiently compute the most probable tag sequence for a given word sequence?
  - Efficiently compute the likelihood for an HMM (probability it outputs a given sequence  $s$ )?
  - Learn the parameters of an HMM given unlabelled training data?
- What are the properties of these algorithms (complexity, convergence, etc)?

# Tagging example

Words:

Possible tags:  
(ordered by  
frequency for  
each word)

<s>	one	dog	bit	</s>
<s>	CD	NN	NN	</s>
	NN	VB	VBD	
	PRP			

# Tagging example

Words:

Possible tags:  
(ordered by  
frequency for  
each word)

<s>	one	dog	bit	</s>
<s>	CD	NN	NN	</s>
	NN	VB	VBD	
	PRP			

- Choosing the best tag for each word independently gives the wrong answer (<s> CD NN NN </s>).
- $P(\text{VBD} | \text{bit}) < P(\text{NN} | \text{bit})$ , but may yield a better *sequence* (<s> CD NN VB </s>)
  - because  $P(\text{VBD} | \text{NN})$  and  $P(\text{</s>} | \text{VBD})$  are high.

# Viterbi: intuition

Words:

Possible tags:  
(ordered by  
frequency for  
each word)

<s>	one	dog	bit	</s>
<s>	CD	NN	NN	</s>
	NN	VB	VBD	
	PRP			

- Suppose we have already computed
  - a) The best tag sequence for <s> ... bit that ends in NN.
  - b) The best tag sequence for <s> ... bit that ends in VBD.
- Then, the best full sequence would be either
  - sequence (a) extended to include </s>, or
  - sequence (b) extended to include </s>.



# Viterbi: intuition

Words:

Possible tags:  
(ordered by  
frequency for  
each word)

<s>	one	dog	bit	</s>
<s>	CD	NN	NN	</s>
	NN	VB	VBD	
	PRP			

- But similarly, to get
  - a) The best tag sequence for <s> ... bit that ends in NN.
- We could extend one of:
  - The best tag sequence for <s> ... dog that ends in NN.
  - The best tag sequence for <s> ... dog that ends in VB.
- And so on...

# Viterbi: high-level picture

- Intuition: the best path of length  $t$  ending in state  $Q$  must include the best path of length  $t-1$  to the previous state (call it  $P$ ). ( $t$  now a *time*, not a *tag*):

$$\langle s \rangle / \langle s \rangle \quad \dots \quad o_{t-1} / P \quad o_t / Q$$

- Because otherwise there must be a better path to  $P$  that we should have used, thereby getting a better path to  $Q$ 
  - Remember the Markov assumptions
    - $P(o_t|Q)$  and  $P(Q|P)$  are *independent* of everything from  $\langle s \rangle$  to  $Q$

# Viterbi: high-level picture

- Intuition: the best path of length  $t$  ending in state  $q$  must include the best path of length  $t-1$  to the previous state. ( $t$  now a *time step*, not a *tag*). So,
  - Find the best path of length  $t-1$  to each state.
  - Consider extending each of those by 1 step, to state  $q$ .
  - Take the best of those options as the best path to state  $q$ .


# Notation

- Sequence of observations over time  $o_1, o_2, \dots, o_T$ 
  - here, words in sentence
- Vocabulary size  $V$  of possible observations
- Set of possible states  $q^1, q^2, \dots, q^N$  (see note next slide)
  - here, tags
- $A$ , an  $N \times N$  matrix of transition probabilities
  - $a_{ij}$ : the prob of transitioning from state  $q^i$  to  $q^j$ . (JM3 Fig 8.7)
- $B$ , an  $N \times V$  matrix of output probabilities
  - $b_i(o_t)$ : the prob of emitting  $o_t$  from state  $q^i$ . (JM3 Fig 8.8)

# Note on notation

- J&M use  $q_1, q_2, \dots, q_N$  for set of states, but *also* use  $q_1, q_2, \dots, q_T$  for state sequence over time.
  - So, just seeing  $q_1$  is ambiguous (though usually disambiguated from context).
  - I'll instead use  $q^i$  for state names, and  $q_t$  for state at time  $t$ .
  - So we could have  $q_t = q^i$ , meaning: the state we're in at time  $t$  is  $q^i$ .

# HMM example w/ new notation



- States  $\{q^1, q^2\}$  (or  $\{<s>, q^1, q^2\}$ )
- Output alphabet  $\{x, y, z\}$

Adapted from Manning & Schuetze, Fig 9.2

# Transition and Output Probabilities

- Transition matrix **A**:

$$a_{ij} = P(q^j | q^i)$$

	$q^1$	$q^2$
$<s>$	1	0
$q^1$	.7	.3
$q^2$	.5	.5

- Output matrix **B**:

$$b_i(o) = P(o | q^i)$$

for output  $o$

	$x$	$y$	$z$
$q^1$	.6	.1	.3
$q^2$	.1	.7	.2

# Joint probability of (states, outputs)

- Let  $\lambda = (A, B)$  be the parameters of our HMM.
- Using our new notation, given state sequence  $Q = (q_1 \dots q_T)$  and output sequence  $O = (o_1 \dots o_T)$ , we have:

$$P(O, Q | \lambda) = \prod_{t=1}^T P(o_t | q_t) P(q_t | q_{t-1})$$



# Joint probability of (states, outputs)

- Let  $\lambda = (A, B)$  be the parameters of our HMM.
- Using our new notation, given state sequence  $Q = (q_1 \dots q_T)$  and output sequence  $O = (o_1 \dots o_T)$ , we have:

$$P(O, Q | \lambda) = \prod_{t=1}^T P(o_t | q_t) P(q_t | q_{t-1})$$

- Or: 
$$P(O, Q | \lambda) = \prod_{t=1}^T b_{q_t}(o_t) a_{q_{t-1}q_t}$$

# Joint probability of (states, outputs)

- Let  $\lambda = (A, B)$  be the parameters of our HMM.
- Using our new notation, given state sequence  $Q = (q_1 \dots q_T)$  and output sequence  $O = (o_1 \dots o_T)$ , we have:

$$P(O, Q|\lambda) = \prod_{t=1}^T P(o_t|q_t)P(q_t|q_{t-1})$$

- Or: 
$$P(O, Q|\lambda) = \prod_{t=1}^T b_{q_t}(o_t) a_{q_{t-1}q_t}$$

- Example:

$$\begin{aligned} P(O = (y, z), Q = (q^1, q^1)|\lambda) &= b_1(y) \cdot b_1(z) \cdot a_{<s>,1} \cdot a_{11} \\ &= (.1)(.3)(1)(.7) \end{aligned}$$

# Viterbi: high-level picture

- Want to find  $\operatorname{argmax}_Q P(Q|O)$
- Intuition: the best path of length  $t$  ending in state  $q$  must include the best path of length  $t-1$  to the previous state. So,
  - Find the best path of length  $t-1$  to each state.
  - Consider extending each of those by 1 step, to state  $q$ .
  - Take the best of those options as the best path to state  $q$ .

# Viterbi algorithm

- Use a **chart** to store partial results as we go
  - $N \times T$  table, where  $v(j,t)$  is the probability\* of the best state sequence for  $o_1 \dots o_t$  that ends in state  $j$ .

\*Specifically,  $v(j,t)$  stores the max of the joint probability  $P(o_1 \dots o_t, q_1 \dots q_{t-1}, q_t = j | \lambda)$

# Viterbi algorithm

- Use a **chart** to store partial results as we go
  - $N \times T$  table, where  $v(j,t)$  is the probability\* of the best state sequence for  $o_1 \dots o_t$  that ends in state  $j$ .
- Fill in columns from left to right, with

$$v(j, t) = \max_{i=1}^N v(i, t-1) \cdot a_{ij} \cdot b_j(o_t)$$

\*Specifically,  $v(j,t)$  stores the max of the joint probability  $P(o_1 \dots o_t, q_1 \dots q_{t-1}, q_t=j | \lambda)$

# Viterbi algorithm

- Use a **chart** to store partial results as we go
  - $N \times T$  table, where  $v(j, t)$  is the probability\* of the best state sequence for  $o_1 \dots o_t$  that ends in state  $j$ .
- Fill in columns from left to right, with
$$v(j, t) = \max_{i=1}^N v(i, t - 1) \cdot a_{ij} \cdot b_j(o_t)$$
- Store a **backtrace** to show, for each cell, which state at  $t-1$  we came from.

\*Specifically,  $v(j, t)$  stores the max of the joint probability  $P(o_1 \dots o_t, q_1 \dots q_{t-1}, q_t = j | \lambda)$

# Example

- Suppose  $O=xzy$ . Our initially empty table:

	$o_1=x$	$o_2=z$	$o_3=y$
$q^1$			
$q^2$			

# Filling the first column


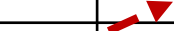
	$o_1=x$	$o_2=z$	$o_3=y$
$q^1$	.6		
$q^2$	0		

$$v(1,1) = a_{<s>1} \cdot b_1(x) = (1)(.6)$$

$$v(2,1) = a_{<s>2} \cdot b_2(x) = (0)(.1)$$



# Starting the second column

	$o_1=x$	$o_2=z$	$o_3=y$
$q^1$	.6		
$q^2$	0		

$$v(1,2) = \max_{i=1}^N v(i,1) \cdot a_{i1} \cdot b_1(z)$$

$$= \max \begin{cases} v(1,1) \cdot a_{11} \cdot b_1(z) = (.6)(.7)(.3) \\ v(2,1) \cdot a_{21} \cdot b_1(z) = (0)(.5)(.3) \end{cases}$$

# Starting the second column

	$o_1=x$	$o_2=z$	$o_3=y$
$q^1$	.6	.126	
$q^2$	0		

$$v(1,2) = \max_{i=1}^N v(i,1) \cdot a_{i1} \cdot b_1(z)$$

$$= \max \begin{cases} v(1,1) \cdot a_{11} \cdot b_1(z) = (.6)(.7)(.3) \\ v(2,1) \cdot a_{21} \cdot b_1(z) = (0)(.5)(.3) \end{cases}$$

# Finishing the second column

	$o_1=x$	$o_2=z$	$o_3=y$
$q^1$	.6	.126	
$q^2$	0		

$$v(2,2) = \max_{i=1}^N v(i,1) \cdot a_{i2} \cdot b_2(z)$$

$$= \max \begin{cases} v(1,1) \cdot a_{12} \cdot b_2(z) = (.6)(.3)(.2) \\ v(2,1) \cdot a_{22} \cdot b_2(z) = (0)(.5)(.2) \end{cases}$$

# Finishing the second column

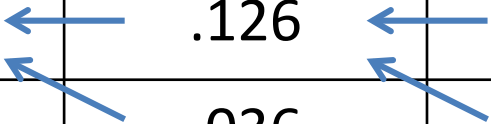
	$o_1=x$	$o_2=z$	$o_3=y$
$q^1$	.6	.126	
$q^2$	0	.036	

$$v(2,2) = \max_{i=1}^N v(i,1) \cdot a_{i2} \cdot b_2(z)$$

$$= \max \begin{cases} v(1,1) \cdot a_{12} \cdot b_2(z) = (.6)(.3)(.2) \\ v(2,1) \cdot a_{22} \cdot b_2(z) = (0)(.5)(.2) \end{cases}$$

# Third column

	$o_1=x$	$o_2=z$	$o_3=y$
$q^1$	.6	.126	.00882
$q^2$	0	.036	.02646



- Exercise: make sure you get the same results!

# Best Path

	$o_1=x$	$o_2=z$	$o_3=y$
$q^1$	.6	.126	.00882
$q^2$	0	.036	.02646

- Choose best final state:  $\max_{i=1}^N v(i, T)$
- Follow backtraces to find best full sequence:  $q^1 q^1 q^2$

# HMMs: what else?

- As with probabilities in N-gram models and classification, chart probabilities get really tiny really fast, risking underflow
  - So, we use **costs** (negative log probabilities) instead
  - Take minimum over sum of costs, instead of maximum over product of probabilities.
- Using Viterbi, we can find the best tags for a sentence (**decoding**), and get  $P(O, Q|\lambda)$ .
- We might also want to
  - Compute the **likelihood**  $P(O|\lambda)$ , i.e., the probability of a sentence regardless of tags (a language model!)
  - **learn** the best set of parameters  $\lambda = (A, B)$  given only an *unannotated* corpus of sentences.

# Computing the likelihood

- From probability theory, we know that

$$P(O|\lambda) = \sum_Q P(O, Q|\lambda)$$

- There are an exponential number of  $Q$ s.
- Again, by computing and storing partial results, we can solve efficiently.
- (Next slides show the algorithm but I'll likely skip them)



# Forward algorithm

- Use a table with cells  $\alpha(j,t)$ : the probability of being in state  $q^t$  after seeing  $o_1 \dots o_t$  (**forward probability**).

$$\alpha(j, t) = P(o_1, o_2, \dots, o_t, q^t = j | \lambda)$$

- Fill in columns from left to right, with

$$\alpha(j, t) = \sum_{i=1}^N \alpha(i, t-1) \cdot a_{ij} \cdot b_j(o_t)$$

- Same as Viterbi, but sum instead of max (and no backtrace).

Note: because there's a sum, we can't use the trick that replaces probabilities with costs. For implementation info, see <http://digital.cs.usu.edu/~cyan/CS7960/hmm-tutorial.pdf> and <http://stackoverflow.com/questions/13391625/underflow-in-forward-algorithm-for-hmms>

# Example

- Suppose  $O=xzy$ . Our initially empty table:

	$o_1=x$	$o_2=z$	$o_3=y$
$q^1$			
$q^2$			

# Filling the first column

	$o_1=x$	$o_2=z$	$o_3=y$
$q^1$	.6		
$q^2$	0		

$$\alpha(1,1) = a_{<s>1} \cdot b_1(x) = (1)(.6)$$

$$\alpha(2,1) = a_{<s>2} \cdot b_2(x) = (0)(.1)$$

# Starting the second column

	$o_1=x$	$o_2=z$	$o_3=y$
$q^1$	.6	.126	
$q^2$	0		

$$\begin{aligned}\alpha(1,2) &= \sum_{i=1}^N \alpha(i,1) \cdot a_{i1} \cdot b_1(z) \\ &= \alpha(1,1) \cdot a_{11} \cdot b_1(z) + \alpha(2,1) \cdot a_{21} \cdot b_1(z) \\ &= (.6)(.7)(.3) + (0)(.5)(.3) \\ &= .126\end{aligned}$$

# Finishing the second column

	$o_1=x$	$o_2=z$	$o_3=y$
$q^1$	.6	.126	
$q^2$	0	.036	

$$\begin{aligned}\alpha(2,2) &= \sum_{i=1}^N \alpha(i,1) \cdot a_{i2} \cdot b_2(z) \\ &= \alpha(1,1) \cdot a_{12} \cdot b_2(z) + \alpha(2,1) \cdot a_{22} \cdot b_2(z) \\ &= (.6)(.3)(.2) + (0)(.5)(.2) \\ &= .036\end{aligned}$$

# Third column and finish

	$o_1=x$	$o_2=z$	$o_3=y$
$q^1$	.6	.126	.01062
$q^2$	0	.036	.03906

- Add up all probabilities in last column to get the probability of the entire sequence:

$$P(O|\lambda) = \sum_{i=1}^N \alpha(i, T)$$

# Learning

- Given *only* the output sequence, learn the best set of parameters  $\lambda = (A, B)$ .
- Assume 'best' = maximum-likelihood.
- Other definitions are possible, won't discuss here.

# Unsupervised learning

- Training an HMM from an annotated corpus is simple.
  - **Supervised** learning: we have examples labelled with the right ‘answers’ (here, tags): no hidden variables in training.
- Training from unannotated corpus is trickier.
  - **Unsupervised** learning: we have no examples labelled with the right ‘answers’: all we see are outputs, state sequence is hidden.



# Circularity

- If we know the state sequence, we can find the best  $\lambda$ .
  - E.g., use MLE:  $P(q^j|q^i) = \frac{c(q^i \rightarrow q^j)}{c(q^i)}$
- If we know  $\lambda$ , we can find the best state sequence.
  - use Viterbi
- But we don't know either!

# Expectation-maximization (EM)

Essentially, a bootstrapping algorithm.

- Initialize parameters  $\lambda^{(0)}$
- At each iteration  $k$ ,
  - E-step: Compute **expected counts** using  $\lambda^{(k-1)}$
  - M-step: Set  $\lambda^{(k)}$  using MLE on the expected counts
- Repeat until  $\lambda$  doesn't change (or other stopping criterion).

# Expected counts??

Counting transitions from  $q^i \rightarrow q^j$ :

- Real counts:
  - count 1 each time we see  $q^i \rightarrow q^j$  in true tag sequence.
- Expected counts:
  - With current  $\lambda$ , compute probs of all possible tag sequences.
  - If sequence  $Q$  has probability  $p$ , count  $p$  for each  $q^i \rightarrow q^j$  in  $Q$ .
  - Add up these fractional counts across all possible sequences.

# Example

- Notionally, we compute expected counts as follows:

Possible sequence				Probability of sequence
$Q_1 =$	$q^1$	$q^1$	$q^1$	$p_1$
$Q_2 =$	$q^1$	$q^2$	$q^1$	$p_2$
$Q_3 =$	$q^1$	$q^1$	$q^2$	$p_3$
$Q_4 =$	$q^1$	$q^2$	$q^2$	$p_4$
Observs:	x	z	y	

# Example

- Notionally, we compute expected counts as follows:

Possible sequence				Probability of sequence
$Q_1 =$	$q^1$	$q^1$	$q^1$	$p_1$
$Q_2 =$	$q^1$	$q^2$	$q^1$	$p_2$
$Q_3 =$	$q^1$	$q^1$	$q^2$	$p_3$
$Q_4 =$	$q^1$	$q^2$	$q^2$	$p_4$
Observs:	x	z	y	

$$\hat{C}(q^1 \rightarrow q^1) = 2p_1 + p_3$$

# Forward-Backward algorithm

- As usual, avoid enumerating all possible sequences.
- **Forward-Backward** (Baum-Welch) algorithm computes expected counts using forward probabilities and **backward probabilities**:


$$\beta(j, t) = P(q^t = j, o_{t+1}, o_{t+2}, \dots o_T | \lambda)$$

– Details, see J&M 6.5

- EM idea is much more general: can use for many latent variable models.


# Guarantees

- EM is guaranteed to find a **local** maximum of the likelihood.



# Guarantees

- EM is guaranteed to find a **local** maximum of the likelihood.



- Not guaranteed to find **global** maximum.
- Practical issues: initialization, random restarts, early stopping.



# Summary

- HMM: a generative model of sentences using hidden state sequence
- Dynamic programming algorithms to compute
  - Best tag sequence given words (Viterbi algorithm)
  - Likelihood (forward algorithm)
  - Best parameters from unannotated corpus (forward-backward algorithm, an instance of EM)